



Team Details

Team Name:

BYTE BRIGADE

SR. NO	ROLE	NAME	ACADEMIC YEAR
1	Team Leader	Prakhar Bajaj	2nd Year
2	Member 1	Raghav Soni	2nd Year
3	Member 2	Parth Kumar	2nd Year
4	Member 3	Mayank Lodhi	2nd Year

 Problem Statement PS01 (KLA) - AI-Based Restoration of Degraded Images for Semiconductor Inspection

 COLLEGE NAME

LNCT University, Bhopal

 TEAM LEADER CONTACT NUMBER

+91 91796 86667

 TEAM LEADER EMAIL ADDRESS

prakharbajaj1551@gmail.com



Problem Statement Addressed



PS01 - KLA: AI-Based Restoration of Degraded Images for Semiconductor Inspection

DESCRIPTION / DETAILS

Wafer inspection images arrive noisy and low-resolution: fast, low-dose scanning trades image quality for throughput.

Three degradations are present - speckle noise (scales with brightness), additive Gaussian sensor noise, and 2x downsampling - applied in an undisclosed order.

Real defects hide inside this noise. A missed defect can cost an entire wafer; a false alarm wastes engineer time.

Task: learn a mapping from a 128x128 degraded image to its clean 256x256 original, scored on PSNR, SSIM, LPIPS and end-to-end speed.


CLEAN 256x256



ground truth

speckle noise
+ Gaussian noise
+ 2x downsample
→
applied in an
undisclosed order

DEGRADED 128x128



what the model receives



Idea Description



One small network restores a noisy 128x128 image to a clean 256x256 one - denoising and 2x super-resolution in a single pass, trained only on the data provided.

KEY CONCEPT & APPROACH

*One small network denoises and 2x super-resolves in a single pass.
Supervised on the 3,200 provided pairs - no external data, no pretrained restoration weights.
Predicts only a correction on a bicubic upscale, which strongly constrains hallucination.*

SOLUTION OVERVIEW

*NAFNet-SR: 0.53 M parameters, 2.3 MB checkpoint - runs on a laptop CPU.
Trained on the judged metrics: Charbonnier + SSIM + LPIPS loss.
28.00 dB PSNR / 0.759 SSIM / 0.175 LPIPS - beats bicubic on all three.
End-to-end 25 ms per image on a T4 GPU (40 images/second).*



Proposed Solution




NAFNet-SR: 24 simplified restoration blocks and a learned 2x upsampler, added to a bicubic upscale so the network predicts only a correction.

SOLUTION DETAILS

WHY NAFNet - SOTA quality, no transformer, no activation function. LOSS - Charbonnier + 0.15 SSIM + 0.05 LPIPS (no GAN loss).

TRAINING - 3,000/100/100 seeded split, 128 px crops, batch 16, 100 epochs, ~2 h on a free Colab T4. INFERENCE - directories in/out, batched, tiled.



the network learns only the **RESIDUAL CORRECTION**
noise to remove and edges to sharpen - not the whole image



Innovation and Uniqueness



We train directly on the judged metrics, and constrain the model to correct rather than generate - decisive when invented structure could be mistaken for a defect.

KEY INNOVATION

We train directly on the scoreboard: SSIM and LPIPS terms sit inside the loss, so the model optimises two of the three judged metrics rather than hoping they follow from pixel accuracy.

Hallucination is strongly constrained: the global bicubic skip means the network predicts a correction rather than generating an image freely.

Float-faithful pipeline: degraded inputs legitimately exceed $[0,1]$ and we never clip them; clipping is applied only to outputs, where ground truth guarantees the range.

COMPETITIVE ADVANTAGE

SMALLER: 0.53 M parameters (2.3 MB) versus 10-26 M for Restormer or SwinIR, so throughput - an explicit scoring axis - is a strength.

FASTER: 25 ms per image end-to-end on a T4; all 400 test images in ~10 s.

TRUSTWORTHY: beats the bicubic baseline on 100 of 100 held-out images; failure cases are published and analysed rather than hidden.

REPRODUCIBLE: seeded splits, one-command training and inference, weights and runtime report committed to the repository.



Impact and Benefits





Better on all three judged image-quality metrics, at 25 ms per image - recovering detail that noise would otherwise hide, with no new hardware.

Primary Impact

Detail hidden inside sensor noise becomes visible again - fewer missed defects and fewer false alarms.
Enables faster or lower-dose scanning: recover quality in software, so existing microscopes inspect more wafers.

Quantifiable Outcomes

PSNR 22.96 -> 28.00 dB (+5.04 dB)
SSIM 0.537 -> 0.759 | LPIPS 0.433 -> 0.175
25 ms/image - all 400 test images in ~10 s
0.53 M parameters, 2.3 MB checkpoint - runs on CPU



grey = bicubic baseline green = ours



Technology & Feasibility/Methodology Used



PyTorch and NAFNet-SR, trained in about two hours on a free Colab T4, and deployable on either a GPU or a plain laptop CPU.

IMPLEMENTATION STRATEGY

*Trained end to end on a free Google Colab T4 in about two hours; no paid infrastructure was used at any stage.
Inference runs on GPU or CPU. Batched on CUDA, tiled for large images, and verified on 128, 256 and non-square inputs.
Deployment is a drop-in folder-to-folder script plus a Streamlit UI for operators; evaluators need no source-code edits.
Measured, not estimated: benchmark.py reports end-to-end runtime with hardware and software versions stated.*



PyTorch 2.11 - NAFNet-SR



NVIDIA T4 / H100 - CPU
capable



Colab - Git - Streamlit



GitHub & Video Link



GitHub Repository

github.com/prakharbajaj1551/semicon-image-restoration



Prototype / Simulation Video

<https://www.youtube.com/watch?v=RMSDaviTOlw>



Research and References



Research Background & Methodology

Every component is drawn from peer-reviewed restoration research, not invented for this hackathon.

Built on published, peer-reviewed research rather than an unproven design: NAFNet (ECCV 2022) matches transformer restorers at a fraction of the cost.

Loss follows restoration literature - Charbonnier for pixel fidelity, SSIM for structure, LPIPS for perceptual quality.

Validated on a held-out split with no selection leakage, against a bicubic baseline, with failure cases analysed.



References & Citations

The three papers behind our architecture, our upsampler and our perceptual metric.

Ref 1: Chen et al., Simple Baselines for Image Restoration (NAFNet), ECCV 2022 - arxiv.org/abs/2204.04676

Ref 2: Zhang et al., The Unreasonable Effectiveness of Deep Features as a Perceptual Metric (LPIPS), CVPR 2018

Ref 3: Shi et al., Real-Time Single Image Super-Resolution Using an Efficient Sub-Pixel CNN, CVPR 2016